

Word-level held-out fit vs frequency

Per-word held-out fit rises with document frequency (one training fit, K = 40)

